

KEVIN CHAN

Salt Lake City, UT | kchan039@gmail.com

Professional Summary

Customer-centric **Forward Deployed Engineer and Enterprise AI Solutions Leader** with 10+ years of experience driving post-sales technical engagements, customer architecture, and cloud-native AI deployments. Expert in guiding high-value accounts from initial discovery and proof-of-concept through full production rollout and organizational enablement. Hands-on practitioner bridging engineering and client environments with custom embeddings, agentic systems, RAG, and model evaluation. Deep technical foundation in HIPAA-compliant healthcare architectures, deployment of complex on-premise or cloud infrastructures, and executive stakeholder advisory.

Core Competencies & Technical Skills

- **Forward Deployment & Customer Success:** Post-Sales Technical Advisory, Solutions Architecture, Production Integrations, Executive Stakeholder Management, Workflow Mapping, Customer Enablement
- **Generative AI & Machine Learning:** OpenAI API Platform, Model Fine-Tuning, Custom Embeddings, RAG Pipelines, Agentic Systems, LangGraph, Model Observability, local deployment frameworks (llama.cpp, Ollama, MLX, OpenVINO)
- **Enterprise Security & Compliance:** HIPAA Compliance, Data Privacy, Healthcare Security Controls, SSO, Identity Management, Middleware Integration
- **Engineering Infrastructure:** AWS, Python, JavaScript/TypeScript, Docker, Kubernetes, CI/CD Pipelines, Terraform, REST APIs, Microservices

Key Professional Experience

Umbrage Studios | Director of Engineering - AI Enablement Sept 2025 – Aug 2026

- Served as the primary deployed technical lead for enterprise GenAI engagements, guiding client teams through workflow mapping, PoC architecture, and production integrations.
- Architected agentic solutions, Model Context Protocol (MCP) servers, and RAG pipelines using OpenAI APIs, while leading hands-on client deployments for model fine-tuning and custom embeddings.
- Partnered directly with C-suite and product stakeholders to map workflows, establish ROI baselines, validate business requirements, and sequence adoption milestones for rapid deployment.
- Conducted technical enablement sessions for client engineering teams to accelerate the adoption of prompt engineering, model observability, and robust deployment patterns.

Accenture | SDE Associate Manager MARCH 2021 – AUG 2025

- Acted as a trusted technical advisor and deployed engineer for Fortune 500 clients across healthcare, scientific research, financial services, and technology sectors.
- Architected and deployed complex cloud solutions adhering to stringent HIPAA regulatory and data privacy requirements directly within client environments.
- Led distributed technical teams of up to 25 engineers across multi-workstream integration initiatives with budgets up to \$2.5M.
- Modernized engineering practices within client ecosystems to reduce delivery timelines by up to 12 months, generating \$500K–\$1M in direct client cost savings and doubling deployment velocity.

Accenture | Software Consultant FEB 2020 – MARCH 2021

- Partnered directly with cross-functional client teams to architect and deliver scalable REST APIs, cloud microservices, and customer-facing interfaces.
- Configured automated CI/CD deployment pipelines on AWS to streamline customer onboarding, integration speed, and account readiness.

Education & Certifications

- **B.S., Computer Engineering** — University of California, Riverside